

Mentoring Operating System (MentOS)

Process management

Created by

Enrico Fraccaroli

enrico.fraccaroli@gmail.com



Table of Contents

1. Process descriptor
 - 1.1. Process identifier
 - 1.2. State of a process
 - 1.3. Time accounting
 - 1.4. Context of a process
2. Scheduler
 - 2.1. Data structures
 - 2.2. Scheduling algorithms
 - 2.3. Context switch



Process descriptor



Process descriptor

The `task_struct` is a data structure used by the Kernel to represent a process and store information about it¹.

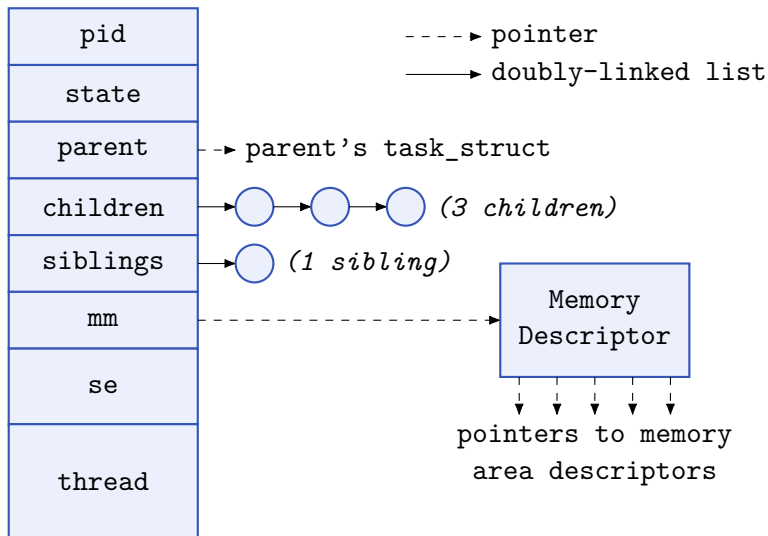
```
struct task_struct {  
    pid_t pid;                // the process identifier  
    unsigned long state;      // the current process's state  
    struct task_struct *parent; // pointer to parent process  
    struct list_head children;  // list of children process  
    struct list_head siblings;  // list of siblings process  
    struct mm_struct *mm;       // memory descriptor  
    struct sched_entity se;      // time accounting (aka schedule entity)  
    struct thread_struct thread; // context of process  
    struct list_head run_list;  // pointer to the process into the scheduler  
}
```

N.B.: The memory descriptor of a process is only reported here for completeness. It will be explained in detail in the Memory Management section.

¹ In Linux, it is quite big, 1.7KB on 32-bit machine (`include/linux/sched.h`)



task_struct memory representation



Process descriptor

Process identifier



Process Identifier (PID) is numeric value identifying a process. When a new process is created a new PID is generated by summing 1 to the last assigned PID.

In Linux, the maximum value for a PID is 32768. When the PID maximum value is reached, the last assigned PID is reset to 0 before searching for a new PID.

The macro `RESERVED_PID` (usually set to 300) is defined to reserved PIDs to system processes and daemons, namely processes proving a service (e.g. a web server). All user's processes have PID greater than `RESERVED_PID`.



Process descriptor

State of a process



State of a process (1/3)

Process state is a numeric value describing the current state of the process. A process can be in one of the following state:

- ▶ **TASK_RUNNING**: either the process is currently in execution, or it has all the resources to be executed except the CPU.
- ▶ **TASK_INTERRUPTIBLE**: the process is blocked (sleep), waiting for some condition to run. When this condition exists, the kernel sets the process's state to **TASK_RUNNING**. The process also awakes and becomes runnable if it receives a signal (e.g., interrupt, signal, released resources).
- ▶ **TASK_UNINTERRUPTIBLE**: this state is identical to **TASK_INTERRUPTIBLE** but it does not depend on specific signal, it must wait without interruption for a specific weak-up call (e.g., task waiting for data transferred from block dev to buffer).



State of a process (2/3)

- ▶ **TASK_STOPPED**: process execution has stopped; the task is not running nor is it eligible to run.
- ▶ **EXIT_ZOMBIE**: Process execution is terminated, but the parent process has not yet issued a `wait4(0)` or `waitpid()` system call to return information about the dead process.
- ▶ **EXIT_DIED**: The final state: the process is being removed by the system because the parent process has just issued a `wait4()` or `waitpid()` system call for it.

Remember init (PID = 1) process.



State of a process (3/3)

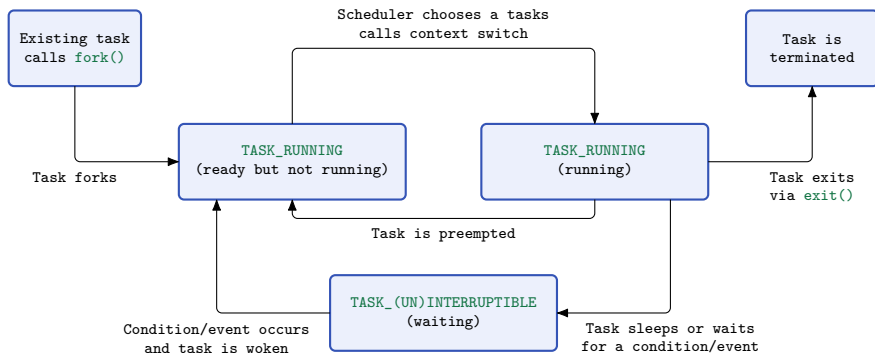


Figure: Flow chart of process states



Relationships among processes (1/2)

Processes created by a program have a parent/child relationship. When a process creates multiple children, these children have sibling relationships.

```
struct task_struct {  
    // ...  
    pid_t pid;           // the process identifier  
    struct task_struct *parent; // pointer to parent process  
    struct list_head children; // list of children process  
    struct list_head siblings; // list of siblings process  
    // ...  
}
```

Fields of `task_struct` describing the relations among processes:

- ▶ **parent**: pointer to the process's parent;
- ▶ **children**: The head of the list containing all children created by the process.
- ▶ **sibling**: The head of the list containing all children created by the process's parent.



Relationships among processes (2/2)




Figure: Parenthood relationships among four processes.

Red lines go from **parent** to **child**.

Blue lines go from **child** to **parent**.

Black lines show relations between **siblings**.



Process descriptor

Time accounting



Time accounting (1/3)

The field **se** of our **task_struct** is a structure called **sched_entity**, which holds all the information about scheduling activities.

```
struct task_struct {  
    //..  
    struct sched_entity se;    // time accounting (aka schedule entity)  
    //..  
}
```

It contains the **priority** and **execution times** of a process.

```
struct sched_entity {  
    int    prio;                // priority  
    time_t start_runtime;       // start execution time  
    time_t exec_start;          // last context switch time  
    time_t sum_exec_runtime;     // overall execution time  
    time_t vruntime;            // weighted execution time  
}
```



Time accounting (2/3)

► **prio**

Defines the execution priority of a process. It has a value in the range [100, 139], where 100 means the highest priority, and 139 means the lowest priority.

By default, the priority of a new generated process is 120.

A process can increment/decrement its **prio** value by using the system call *nice(inc)*, which takes as input parameter a value in the range [-20, 19].

Examples:

- *nice(1)* (increment **prio** value of calling process by 1 unit)
120 \Rightarrow 121
- *nice(-5)* (decrement **prio** value of calling process by -5 units)
120 \Rightarrow 115



Time accounting (3/3)

- ▶ **start_runtime**

The system execution time reporting when the process was first executed in the CPU.

- ▶ **exec_start**

The system execution time reporting when the process was last executed in the CPU.

- ▶ **sum_exec_runtime**

The overall execution time spent by the process in CPU.

- ▶ **vruntime**

The virtual runtime, namely the weighted overall execution time spent by the process in CPU (see CFS).



Process descriptor

Context of a process



Context of a process

The field **thread** of our `task_struct` is a structure called `thread_struct`, which holds all the information about the execution of a process.

```
struct task_struct {  
    //..  
    struct thread_struct thread; // context of process  
}
```

It is called the **context** of a process, and whenever a process is **not running**, it contains all the vital information required to **resume** it.

```
struct thread_struct {  
    uint32_t ebp;        // base pointer register  
    uint32_t esp;        // stack pointer register  
    uint32_t ebx;        // base register  
    uint32_t edx;        // data register  
    uint32_t ecx;        // counter  
    uint32_t eax;        // accumulator register  
    uint32_t eip;        // Instruction Pointer Register  
    uint32_t eflags;      // flag register  
    bool_t fpu_enabled;   // is FPU enabled?  
    savefpu fpu_register; // FPU context  
}
```



Scheduler



Scheduler

Data structures



Scheduler data structures

The **runqueue** data structure is the most important data structure of the scheduler. It collects all system processes in running state.

```
struct runqueue {  
    unsigned long nr_running; // number of processes in running state  
    struct task_struct *curr; // pointer to current running process  
    struct list_head_t queue; // list of processes in running state  
}
```

Pay attention!

The **queue** field is the *head* of a circular, doubly-linked list collecting all system processes in running state. Consequently, a field **run_list** of type *struct list_head* is added in the *struct task_struct*.

(see slides fundamental concepts for more details).



Scheduler execution flow

The scheduler is called after the handle of an interrupt/exception. In detail, the following operations are performed by the scheduler:

1. updates the time accounting variables of the current process;
2. tries to wake up a waiting process. Whether a waiting condition is met, a process is woken by setting its state to running, and inserting it into the runqueue (topic not faced in current slides);
3. run scheduling algorithm to pick the next process to be executed by CPU from the runqueue;
4. performs context switch.



Scheduler

Scheduling algorithms



Select next process

`scheduler_pick_next_task` is the function used by the scheduler to get the next process to execute, and internally this function calls the currently selected scheduling algorithm.

Based on the selected scheduling algorithm, the next process can be chosen differently.

MentOS provides the following tree algorithms:

- ▶ **RR** - Round-Robin (`__scheduler_rr`);
- ▶ **Priority** - Highest Priority First (`__scheduler_priority`);
- ▶ **CFS** - Completely Fair Scheduler (`__scheduler_cfs`).

Pay attention!

In the following algorithms, we use the doubly-linked list defined in Linux Kernel, to collect all processes in running state.



Select next process (Round-Robin) (1/4)

Round Robin is a CPU scheduling algorithm where a fixed time slice is assigned to each system process, in a cyclic way. It is simple, preemptive, easy to implement, and starvation-free.

Pseudocode of Round-Robin algorithm.

Require: Current process c , List of processes L

Ensure: Next process n

- 1: $\text{nextNode} = \text{next}(c)$
- 2: **if** $\text{IsTheHead}(L, \text{nextNode})$ **then**
- 3: $\text{nextNode} = \text{next}(\text{nextNode})$
- 4: **end if**
- 5: $n = \text{list_entry}(\text{nextNode})$



Select next process (Round-Robin) (2/4)

Now, given the runqueue:

```
struct runqueue {
    unsigned long nr_running; // number of processes in running state
    struct task_struct *curr; // pointer to current running process
    struct list_head_t queue; // list of processes in running state
}
```

here is a **simplified** implementation of the Round-Robin algorithm:

```
struct task_struct * __scheduler_rr(struct runqueue *runqueue) {
    // nextNode = next(c)
    struct list_head *nextNode = runqueue->curr->run_list.next;

    // if IsTheHead(L, nextNode)
    if (nextNode == &runqueue->queue)
        nextNode = nextNode->next;

    // n = entry(nextNode)
    task_struct *n = list_entry(nextNode, struct task_struct, run_list);
    return n;
}
```



Select next process (Round-Robin) (3/4)

The actual implementation in MentOs considers the presence of **periodic** processes, which are discussed in the **Real-Time Scheduler** slides. As such, it is slightly more complex than what is shown in the previous slide.

However, the idea stays the same, except we need to use a **for** loop to **search** for a viable next process. We need a **for** loop because the **next** process might be a **periodic** process, and we might want to **skip** it.



Select next process (Round-Robin) (4/4)

The final code of our Round-Robin will look like this:

```
static inline task_struct *__scheduler_rr(runqueue_t *runqueue, bool_t skip_periodic)
{
    // If there is just one task, return it; no need to do anything.
    if (list_head_size(&runqueue->curr->run_list) <= 1) {
        return runqueue->curr;
    }
    // Search for the next task (we do not start from the head, so INSIDE, skip the head).
    list_for_each_decl(it, &runqueue->curr->run_list)
    {
        // Check if we reached the head of list_head, and skip it.
        if (it == &runqueue->queue)
            continue;
        // Get the current entry.
        task_struct *entry = list_entry(it, task_struct, run_list);
        // We consider only runnable processes
        if (entry->state != TASK_RUNNING)
            continue;
        // If entry is a periodic task, and we were asked to skip periodic tasks, skip it.
        if (__is_periodic_task(entry) && skip_periodic)
            continue;
        // We have our next entry.
        return entry;
    }
    return NULL;
}
```




Example (Round-Robin) (1/7)

First iteration:

- ▶ `current_process = init`
- ▶ `__scheduler_rr()` returns `shell`

runqueue:




Example (Round-Robin) (2/7)

Second iteration:

- ▶ `current_process = shell`
- ▶ `__scheduler_rr()` returns `task1`

runqueue:




Example (Round-Robin) (3/7)

Third iteration:

- ▶ `current_process = task1`
- ▶ `__scheduler_rr()` returns `task2`

runqueue:




Example (Round-Robin) (4/7)

Fourth iteration:

- ▶ `current_process = task2`
- ▶ `__scheduler_rr()` returns `task3`

runqueue:




Example (Round-Robin) (5/7)

Fifth iteration:

- ▶ `current_process = task3`
- ▶ `__scheduler_rr()` returns `task4`

runqueue:




Example (Round-Robin) (6/7)

Sixth iteration:

- ▶ `current_process = task4`
- ▶ `__scheduler_rr()` returns `init`

runqueue:




Example (Round-Robin) (7/7)

Seventh iteration:

- ▶ `current_process = init`
- ▶ `__scheduler_rr()` returns `shell`

runqueue:



Select next process (Highest Priority First) (1/3)

Round robin scheduling assumes that all processes are equally important. This generally is untrue. We would sometimes like to see long CPU-intensive (non-interactive) processes get a lower priority than interactive processes.

In addition, different users may have different status. A system administrator's processes may rank above those of a student's.

These goals led to the introduction of the *Priority* scheduling algorithm.



Select next process (Highest Priority First) (2/3)

Each process has a static priority. Smaller is the number, higher is the priority of the process.

The scheduler simply picks the highest priority process to run. A process is **preempted** whenever a higher priority process is available in the run queue.

Advantage: priority scheduling provides a good mechanism where the relative importance of each process may be precisely defined.

Disadvantage: If high priority processes use up a lot of CPU time, lower priority processes may starve and be postponed indefinitely, leading to **starvation**.



Select next process (Highest Priority First) (3/3)

Pseudocode of Highest Priority First.

Require: Current process c , List of processes L

Ensure: Next process n

```
1:  $n = c$ 
2: for all  $listNode \in L$  do
3:   if !IsTheHead( $L$ ,  $listNode$ ) then
4:      $t = list\_entry(listNode)$ 
5:     if  $priority(t) < priority(n)$  then
6:        $n = t$ 
7:     end if
8:   end if
9: end for
10: return  $n$ 
```

The implementation of this algorithm is given to the student.




Example (Highest Priority First) (1/3)

First block of iteration:

- ▶ `current_process = task2`
- ▶ `__scheduler_priority()` returns `task2` until no process with an higher priority is present in the system.

runqueue:




Example (Highest Priority First) (2/3)

Second block iteration:

- ▶ `current_process = task1`
- ▶ `__scheduler_priority()` returns `task1` until no process with an higher priority is present in the system.

runqueue:



Example (Highest Priority First) (3/3)

Third block of iteration:

- ▶ `current_process = task3`
- ▶ `__scheduler_priority()` returns `task3` until no process with an higher priority is present in the system.

runqueue:

<code>init</code> prio = 120	<code>shell</code> prio = 120	<code>task1</code> prio = 110	<code>task3</code> prio = 105
---------------------------------	----------------------------------	----------------------------------	----------------------------------



How much time do `init` and `shell` have to wait to get the CPU?



Select next process (Completely Fair Scheduler) (1/6)

Completely Fair Scheduler (CFS) aims to prevent starvation by assigning the CPU fairly to all system processes.

Let consider an example to illustrate the goal of CFS. If there are two tasks A and B, which have a same "weight", the portion of available CPU time given to each task is 50%.

However, if the "weight" of task A increases on CPU by 10%, then task A's portion of the CPU is 55%, meanwhile task B's portion of the CPU becomes 45%.



Select next process (Completely Fair Scheduler) (2/6)

CFS's idea: let use the priority of each process to "weight" its overall execution time (virtual runtime).

Processes having a low priority have a virtual runtime increasing faster than processes with a higher priority. Scheduler always picks the process with the lowest virtual execution time!



Select next process (Completely Fair Scheduler) (3/6)

Scheduler needs to know the weight of the task to estimate its CPU time's portion. Hence, the priority number has to be mapped to such a weight; this is done in the array `prio_to_weight`:

```
static const int prio_to_weight[] = {  
    /* 100 */    88761, 71755, 56483, 46273, 36291,  
    /* 105 */    29154, 23254, 18705, 14949, 11916,  
    /* 110 */    9548, 7620, 6100, 4904, 3906,  
    /* 115 */    3121, 2501, 1991, 1586, 1277,  
    /* 120 */    1024, 820, 655, 526, 423,  
    /* 125 */    335, 272, 215, 172, 137,  
    /* 130 */    110, 87, 70, 56, 45,  
    /* 135 */    36, 29, 23, 18, 15  
};
```



Select next process (Completely Fair Scheduler) (4/6)

A priority number of 120, which is the priority of a normal task, is mapped to a weight of 1024.

Note that the ratio of two successive entries in the array is almost 1.25. This number is chosen such that:

- ▶ if the priority of a task is reduced by one, then it gets 10% higher share of the available CPU time.
- ▶ if the priority of a task is increased by one, then it gets 10% lower share of the available CPU time.



Select next process (Completely Fair Scheduler) (5/6)

Given the array `prio_to_weight` we can update the virtual runtime of a process `p`, namely its weighted overall execution by using the formula:

```
vruntime += delta_exec * (NICE_0_LOAD / weight(p))
```

where:

- ▶ `vruntime` is the virtual run time of the process;
- ▶ `delta_exec` is the last amount of time spent by `p` in the CPU;
- ▶ `NICE_0_LOAD` is the weight of a task with normal priority (1024);
- ▶ `weight(p)` is the weight of `p` defined by the array `prio_to_weight`.



Select next process (Completely Fair Scheduler) (6/6)

Pseudocode of Completely Fair Scheduler.

Require: Current process c , List of processes L

Ensure: Next process n

```
1: updateVirtualRuntime(c)
2:  $n = c$ 
3: for all  $listNode \in L$  do
4:   if !IsTheHead( $L$ ,  $listNode$ ) then
5:      $task = list\_entry(listNode)$ 
6:     if  $virtualRuntime(task) < virtualRuntime(n)$  then
7:        $n = task$ 
8:     end if
9:   end if
10: end for
11: return  $n$ 
```

The implementation of this algorithm is given to the students.

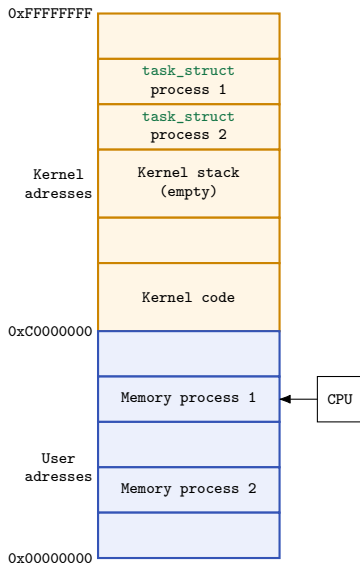


Scheduler

Context switch



Context switch

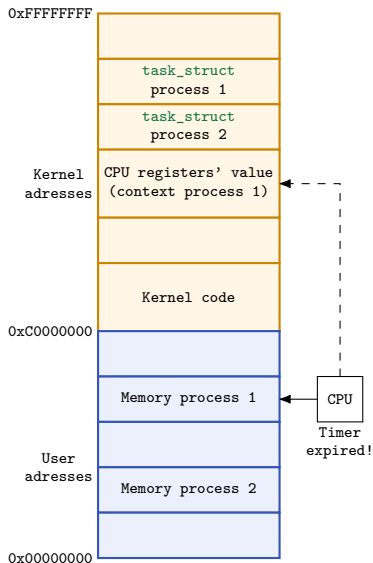


The CPU performs a context switch to change the process executed by CPU.

The following example shows the steps performed by the operating system to save the current process's state (*process 1*), and then resume the execution of a previously stopped process (*process 2*).



Context switch




(1) Time expired! It is time to give back control of CPU to kernel.
Timer device rises the signal *INTR* and present 0 in *irq* line.

When *INTR* is risen, the CPU moves from Ring 3 (user mode) to Ring 0 (kernel mode). After the CPU privilege level change, the values of CPU registers are pushed in the Kernel's stack.

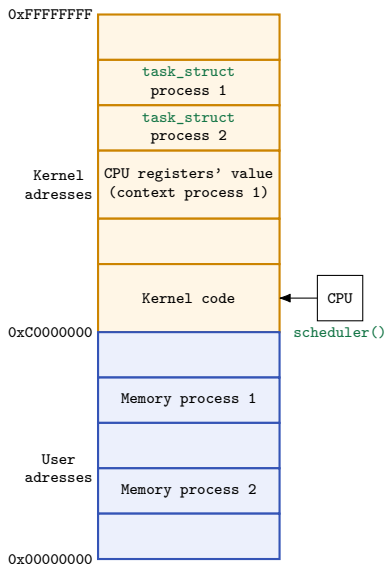


Context switch

(2) CPU starts executing *irq_0* interrupt handler to handle the hardware interrupt 0 risen by Timer.



Context switch




(3) The scheduler is then called to update the time accounting variables of the interrupted process, and pick the next process to run.

In this example, the scheduler picks the process 2 as the next one.



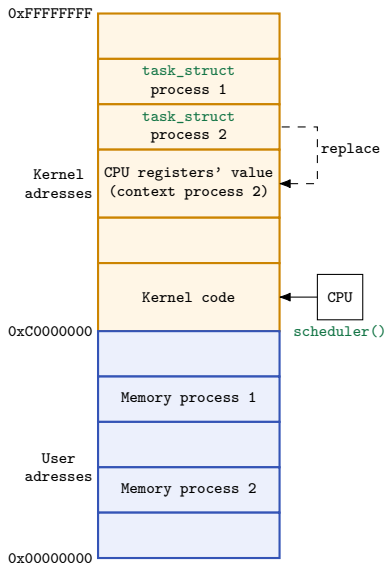
Context switch



(4) Kernel updates the `thread_struct` structure of the `task_struct` of the process 1 in order to save its context.



Context switch

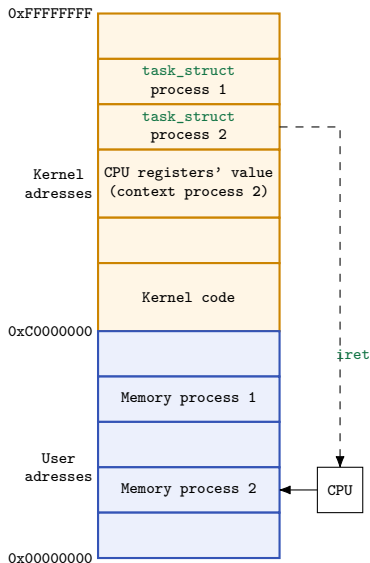


(4) Kernel updates the `thread_struct` structure of the `task_struct` of the process 1 in order to save its context.

(5) Kernel replaces the context of process 1 with the context of process 2 in its stack memory.



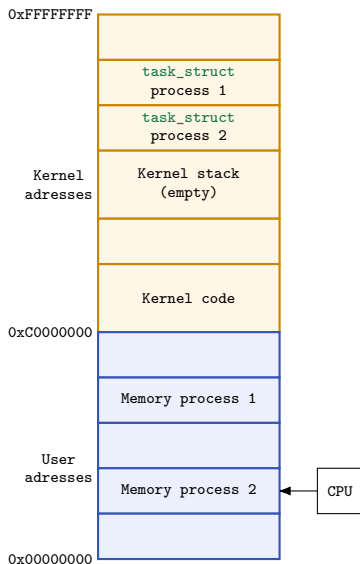
Context switch



(6) Kernel moves the values from its stack to CPU's registers and runs an *iret* assembly instruction, which changes the CPU privilege level from Ring 0 (kernel mode) to Ring 3 (user mode).



Context switch



(7) The context of the process 2 is in the CPU's registers finally. The CPU can keep on executing the code of the process 2 in user mode until the next context switch.

